

Imagenet

December 14, 2025

1 CS 441 final project

```
[3]: import os
import numpy as np
from PIL import Image
import time
import torch
from torch.utils.data import Dataset, DataLoader, random_split
from torchvision import transforms, models

from sklearn.metrics import classification_report, confusion_matrix
from sklearn.svm import LinearSVC
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import SGDClassifier
from sklearn.decomposition import PCA

import torch.nn as nn
import torch.optim as optim
DRIVE_BASE = "/content/drive/My Drive/CS441/Final/garbage_dataset"
LOCAL_TRAIN = "/content/garbage_train"
LOCAL_TEST = "/content/garbage_test"

FORCE_SYNC = False

from google.colab import drive
drive.mount("/content/drive")

!pip -q install pillow-heif
from PIL import Image
import pillow_heif
pillow_heif.register_heif_opener()

import os, subprocess

IMG_EXTS = (".jpg", ".jpeg", ".png", ".bmp", ".webp", ".heic", ".heif")

def count_images(root):
```

```

if not os.path.exists(root):
    return 0
c = 0
for _, _, files in os.walk(root):
    for f in files:
        if f.lower().endswith(IMG_EXTS):
            c += 1
return c

def rsync_dir(src, dst):
    cmd = ["rsync", "-a", "--delete", "--info=progress2", src.rstrip("/") + "/"
        , dst.rstrip("/") + "/"]
    print("Running:", " ".join([f'"{x}"' if " " in x else x for x in cmd]))
    subprocess.run(cmd, check=True)

drive_train = os.path.join(DRIVE_BASE, "train")
drive_test = os.path.join(DRIVE_BASE, "test")

if not os.path.exists(drive_train):
    raise FileNotFoundError(f"Not found: {drive_train}")
if not os.path.exists(drive_test):
    raise FileNotFoundError(f"Not found: {drive_test}")

drive_train_cnt = count_images(drive_train)
drive_test_cnt = count_images(drive_test)

local_train_cnt = count_images(LOCAL_TRAIN)
local_test_cnt = count_images(LOCAL_TEST)

print(f"Drive train images: {drive_train_cnt} | Local train images: {local_train_cnt}")
print(f"Drive test images: {drive_test_cnt} | Local test images: {local_test_cnt}")

need_sync = FORCE_SYNC or (drive_train_cnt != local_train_cnt) or
            (drive_test_cnt != local_test_cnt)
print("Need sync:", need_sync)

if need_sync:
    rsync_dir(drive_train, LOCAL_TRAIN)
    rsync_dir(drive_test, LOCAL_TEST)

train_root = LOCAL_TRAIN
test_root = LOCAL_TEST

local_train_cnt2 = count_images(train_root)
local_test_cnt2 = count_images(test_root)

```

```

print("\n==== READY ===")
print("train_root =", train_root, "| images:", local_train_cnt2)
print("test_root  =", test_root,  "| images:", local_test_cnt2)

```

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).
Drive train images: 20371 | Local train images: 15943
Drive test  images: 251  | Local test  images: 0
Need sync: True
Running: rsync -a --delete --info=progress2 "/content/drive/My
Drive/CS441/Final/garbage_dataset/train/" /content/garbage_train/
Running: rsync -a --delete --info=progress2 "/content/drive/My
Drive/CS441/Final/garbage_dataset/test/" /content/garbage_test/

==== READY ===
train_root = /content/garbage_train | images: 20371
test_root  = /content/garbage_test | images: 251

```

[4]:

```

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print("Device:", device)

```

Device: cuda

[5]:

```

class GarbageFineDataset(Dataset):
    def __init__(self, root_dir, transform=None,
                 fine_class_to_idx=None, big_class_to_idx=None,
                 strict=False,
                 extensions=('.jpg', '.jpeg', '.png', '.bmp', '.webp', '.heic', u
                +'heif')):
        self.root_dir = root_dir
        self.transform = transform
        self.extensions = extensions
        self.strict = strict

        self.fine_class_to_idx = {} if fine_class_to_idx is None else u
        ↵dict(fine_class_to_idx)
        self.big_class_to_idx = {} if big_class_to_idx is None else u
        ↵dict(big_class_to_idx)

        self.samples = [] # (image_path, fine_idx, big_idx)

        big_names = sorted([d for d in os.listdir(root_dir) if os.path.isdir(os.
        ↵path.join(root_dir, d))])
        for big_name in big_names:
            big_path = os.path.join(root_dir, big_name)

```

```

        if big_name in self.big_class_to_idx:
            big_idx = self.big_class_to_idx[big_name]
        else:
            if self.strict:
                raise ValueError(f"[STRICT] test/show train/not exist big_{category}: {big_name}")
            big_idx = len(self.big_class_to_idx)
            self.big_class_to_idx[big_name] = big_idx

        fine_names = sorted([d for d in os.listdir(big_path) if os.path.isdir(os.path.join(big_path, d))])
        for fine_name in fine_names:
            fine_path = os.path.join(big_path, fine_name)

            fine_full_name = f"{big_name}/{fine_name}"

            if fine_full_name in self.fine_class_to_idx:
                fine_idx = self.fine_class_to_idx[fine_full_name]
            else:
                if self.strict:
                    raise ValueError(f"[STRICT] test/show train/not exist small category: {fine_full_name}")
                fine_idx = len(self.fine_class_to_idx)
                self.fine_class_to_idx[fine_full_name] = fine_idx

            for fname in os.listdir(fine_path):
                if fname.lower().endswith(self.extensions):
                    img_path = os.path.join(fine_path, fname)
                    self.samples.append((img_path, fine_idx, big_idx))

        print(f"[{root_dir}] samples={len(self.samples)}, big={len(self.big_class_to_idx)}, fine={len(self.fine_class_to_idx)}")

    def __len__(self):
        return len(self.samples)

    def __getitem__(self, idx):
        img_path, fine_idx, big_idx = self.samples[idx]
        img = Image.open(img_path).convert("RGB")
        if self.transform:
            img = self.transform(img)
        return img, fine_idx, big_idx

    @property
    def fine_idx_to_name(self):
        return {v: k for k, v in self.fine_class_to_idx.items()}

```

```

@property
def big_idx_to_name(self):
    return {v: k for k, v in self.big_class_to_idx.items()}

@property
def fine_to_big(self):
    # fine_idx -> big_idx
    mapping = {}
    for _, f_idx, b_idx in self.samples:
        mapping[f_idx] = b_idx
    return mapping

```

```

[6]: img_size = 224
batch_size = 32
num_workers = 4

train_transform = transforms.Compose([
    transforms.Resize((256, 256)),
    transforms.RandomResizedCrop(img_size),
    transforms.RandomHorizontalFlip(),
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]),
])

val_test_transform = transforms.Compose([
    transforms.Resize((img_size, img_size)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]),
])

# 0.9 / 0.1
full_train_dataset = GarbageFineDataset(train_root, transform=train_transform)

full_train_for_val = GarbageFineDataset(
    train_root,
    transform=val_test_transform,
    fine_class_to_idx=full_train_dataset.fine_class_to_idx,
    big_class_to_idx=full_train_dataset.big_class_to_idx,
    strict=True
)

seed = 42
g = torch.Generator().manual_seed(seed)

```

```

N = len(full_train_dataset)
n_train = int(0.9 * N)
n_val = N - n_train
train_subset, val_subset = random_split(range(N), [n_train, n_val], generator=g)

class IndexSubset(Dataset):
    def __init__(self, base_dataset, indices):
        self.base = base_dataset
        self.indices = list(indices)
    def __len__(self):
        return len(self.indices)
    def __getitem__(self, i):
        return self.base[self.indices[i]]

train_dataset = IndexSubset(full_train_dataset, train_subset)
val_dataset = IndexSubset(full_train_for_val, val_subset)

test_dataset = GarbageFineDataset(
    test_root,
    transform=val_test_transform,
    fine_class_to_idx=full_train_dataset.fine_class_to_idx,
    big_class_to_idx=full_train_dataset.big_class_to_idx,
    strict=True
)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True, □
    ↵num_workers=num_workers)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False, □
    ↵num_workers=num_workers)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False, □
    ↵num_workers=num_workers)

fine_idx_to_name = full_train_dataset.fine_idx_to_name
big_idx_to_name = full_train_dataset.big_idx_to_name
fine_to_big = full_train_dataset.fine_to_big

print("Big classes:", big_idx_to_name)
print("Fine classes (examples):", list(fine_idx_to_name.items())[:11])

```

```

[/content/garbage_train] samples=20371, big=4, fine=16
[/content/garbage_train] samples=20371, big=4, fine=16
[/content/garbage_test] samples=251, big=4, fine=16
Big classes: {0: 'recycling', 1: 'special', 2: 'trash', 3: 'yard_waste'}
Fine classes (examples): [(0, 'recycling/cardboard'), (1, 'recycling/glass'),
(2, 'recycling/metal'), (3, 'recycling/paper'), (4, 'recycling/plastic'), (5,

```

```
'special/battery'), (6, 'special/cables'), (7, 'special/keyboard'), (8,  
'special/mouse'), (9, 'special/tire'), (10, 'trash/biological')]
```

2 1. SVM

```
[7]: svm_transform = transforms.Compose([  
    transforms.Resize((96, 96)),  
    transforms.ToTensor(),  
)  
svm_full_train = GarbageFineDataset(  
    train_root,  
    transform=svm_transform,  
    fine_class_to_idx=full_train_dataset.fine_class_to_idx,  
    big_class_to_idx=full_train_dataset.big_class_to_idx,  
    strict=True  
)  
svm_train_dataset = IndexSubset(svm_full_train, train_subset)  
svm_val_dataset = IndexSubset(svm_full_train, val_subset)  
  
svm_test_dataset = GarbageFineDataset(  
    test_root,  
    transform=svm_transform,  
    fine_class_to_idx=full_train_dataset.fine_class_to_idx,  
    big_class_to_idx=full_train_dataset.big_class_to_idx,  
    strict=True  
)  
  
svm_batch_size = 64  
svm_train_loader = DataLoader(svm_train_dataset, batch_size=svm_batch_size, □  
    ↪shuffle=False, num_workers=num_workers)  
svm_val_loader = DataLoader(svm_val_dataset, batch_size=svm_batch_size, □  
    ↪shuffle=False, num_workers=num_workers)  
svm_test_loader = DataLoader(svm_test_dataset, batch_size=svm_batch_size, □  
    ↪shuffle=False, num_workers=num_workers)  
  
@torch.no_grad()  
def extract_flat_features(loader):  
    feats, labels = [], []  
    for images, fine_labels, _ in loader:  
        flat = images.view(images.size(0), -1)  
        feats.append(flat.cpu().numpy())  
        labels.append(fine_labels.numpy())  
    return np.concatenate(feats), np.concatenate(labels)
```

```

svm_train_feats, svm_train_labels = extract_flat_features(svm_train_loader)
svm_val_feats, svm_val_labels = extract_flat_features(svm_val_loader)
svm_test_feats, svm_test_labels = extract_flat_features(svm_test_loader)

scaler = StandardScaler()
pca_components = 256
svm_train_feats = scaler.fit_transform(svm_train_feats)
svm_val_feats = scaler.transform(svm_val_feats)
svm_test_feats = scaler.transform(svm_test_feats)

pca = PCA(n_components=pca_components, random_state=42)
svm_train_feats = pca.fit_transform(svm_train_feats)
svm_val_feats = pca.transform(svm_val_feats)
svm_test_feats = pca.transform(svm_test_feats)

svm_clf = SGDClassifier(
    loss="hinge",
    alpha=1e-5,
    max_iter=500,
    tol=1e-4,
    n_jobs=-1,
    class_weight="balanced",
    random_state=42,
    verbose=0,
)
t0 = time.time()
svm_clf.fit(svm_train_feats, svm_train_labels)
svm_val_pred = svm_clf.predict(svm_val_feats)
svm_test_pred = svm_clf.predict(svm_test_feats)

svm_val_acc = (svm_val_pred == svm_val_labels).mean()
svm_test_acc = (svm_test_pred == svm_test_labels).mean()

print(f"[SVM] Val acc={svm_val_acc:.4f} | Test acc={svm_test_acc:.4f}")

```

```

[/content/garbage_train] samples=20371, big=4, fine=16
[/content/garbage_test] samples=251, big=4, fine=16
[SVM] Val acc=0.3734 | Test acc=0.1195

```

3 2. CNN

```

[8]: num_fine_classes = len(full_train_dataset.fine_class_to_idx)
class SimpleCNN(nn.Module):
    def __init__(self, num_classes):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),

```

```

        nn.BatchNorm2d(32),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2),
        nn.Conv2d(32, 64, kernel_size=3, padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(2),
        nn.Conv2d(64, 128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(inplace=True),
        nn.AdaptiveAvgPool2d((1, 1)),
    )
    self.classifier = nn.Linear(128, num_classes)

def forward(self, x):
    x = self.features(x)
    x = torch.flatten(x, 1)
    return self.classifier(x)

def train_cnn_one_epoch(model, loader, optimizer, criterion, device):
    model.train()
    total_loss, correct, total = 0.0, 0, 0
    for images, fine_labels, _ in loader:
        images = images.to(device)
        fine_labels = fine_labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, fine_labels)
        loss.backward()
        optimizer.step()

        total_loss += loss.item() * images.size(0)
        pred = outputs.argmax(dim=1)
        correct += (pred == fine_labels).sum().item()
        total += fine_labels.size(0)
    return total_loss / total, correct / total

@torch.no_grad()
def eval_cnn_one_epoch(model, loader, criterion, device):
    model.eval()
    total_loss, correct, total = 0.0, 0, 0
    for images, fine_labels, _ in loader:
        images = images.to(device)
        fine_labels = fine_labels.to(device)

```

```

outputs = model(images)
loss = criterion(outputs, fine_labels)

total_loss += loss.item() * images.size(0)
pred = outputs.argmax(dim=1)
correct += (pred == fine_labels).sum().item()
total += fine_labels.size(0)
return total_loss / total, correct / total

cnn_model = SimpleCNN(num_fine_classes).to(device)
cnn_criterion = nn.CrossEntropyLoss()
cnn_optimizer = optim.Adam(cnn_model.parameters(), lr=1e-3)
cnn_epochs = 10
cnn_best_acc = 0.0
cnn_best_state = None

for epoch in range(cnn_epochs):
    tr_loss, tr_acc = train_cnn_one_epoch(cnn_model, train_loader, □
    ↵cnn_optimizer, cnn_criterion, device)
    va_loss, va_acc = eval_cnn_one_epoch(cnn_model, val_loader, cnn_criterion, □
    ↵device)
    print(f"[CNN] Epoch {epoch+1}/{cnn_epochs} | Train loss={tr_loss:.4f}, □
    ↵acc={tr_acc:.4f} | Val loss={va_loss:.4f}, acc={va_acc:.4f}")
    if va_acc > cnn_best_acc:
        cnn_best_acc = va_acc
        cnn_best_state = {k: v.cpu() for k, v in cnn_model.state_dict().items()}

if cnn_best_acc > 0:
    cnn_model.load_state_dict(cnn_best_state)
    cnn_model = cnn_model.to(device)
print("[CNN] Best val acc:", cnn_best_acc)

```

[CNN] Epoch 1/10 | Train loss=1.8990, acc=0.3824 | Val loss=1.7342, acc=0.4642
[CNN] Epoch 2/10 | Train loss=1.7634, acc=0.4262 | Val loss=1.6741, acc=0.4769
[CNN] Epoch 3/10 | Train loss=1.7054, acc=0.4435 | Val loss=1.6038, acc=0.5020
[CNN] Epoch 4/10 | Train loss=1.6543, acc=0.4612 | Val loss=1.5930, acc=0.5079
[CNN] Epoch 5/10 | Train loss=1.6063, acc=0.4768 | Val loss=1.6048, acc=0.4902
[CNN] Epoch 6/10 | Train loss=1.5679, acc=0.4903 | Val loss=1.4820, acc=0.5206
[CNN] Epoch 7/10 | Train loss=1.5332, acc=0.5009 | Val loss=1.4561, acc=0.5402
[CNN] Epoch 8/10 | Train loss=1.5196, acc=0.5062 | Val loss=1.3990, acc=0.5559
[CNN] Epoch 9/10 | Train loss=1.4680, acc=0.5202 | Val loss=1.4069, acc=0.5648
[CNN] Epoch 10/10 | Train loss=1.4526, acc=0.5232 | Val loss=1.2763, acc=0.5942
[CNN] Best val acc: 0.5942100098135427

4 3. Fine tuned Resnet 18

Stage 1: Freeze the backbone and only train the last layer

```
[9]: weights = models.ResNet18_Weights.IMGNET1K_V1
model = models.resnet18(weights=weights)

for p in model.parameters():
    p.requires_grad = False

num_features = model.fc.in_features
num_fine_classes = len(full_train_dataset.fine_class_to_idx)
model.fc = nn.Linear(num_features, num_fine_classes)

model = model.to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters(), lr=1e-3)
```

Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to
/root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth

100% | 44.7M/44.7M [00:00<00:00, 97.2MB/s]

```
[10]: def train_one_epoch(model, loader, optimizer, criterion, device):
        model.train()
        total_loss, correct, total = 0.0, 0, 0

        for images, fine_labels, _ in loader:
            images = images.to(device)
            fine_labels = fine_labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, fine_labels)
            loss.backward()
            optimizer.step()

            total_loss += loss.item() * images.size(0)
            pred = outputs.argmax(dim=1)
            correct += (pred == fine_labels).sum().item()
            total += fine_labels.size(0)

        return total_loss / total, correct / total

@torch.no_grad()
def eval_one_epoch(model, loader, criterion, device):
    model.eval()
```

```

total_loss, correct, total = 0.0, 0, 0

for images, fine_labels, _ in loader:
    images = images.to(device)
    fine_labels = fine_labels.to(device)

    outputs = model(images)
    loss = criterion(outputs, fine_labels)

    total_loss += loss.item() * images.size(0)
    pred = outputs.argmax(dim=1)
    correct += (pred == fine_labels).sum().item()
    total += fine_labels.size(0)

return total_loss / total, correct / total

```

```

[12]: num_epochs = 10
best_val_acc = 0.0
best_state = None

for epoch in range(num_epochs):
    tr_loss, tr_acc = train_one_epoch(model, train_loader, optimizer, criterion, device)
    va_loss, va_acc = eval_one_epoch(model, val_loader, criterion, device)

    print(f"Epoch {epoch+1}/{num_epochs} | "
          f"Train loss={tr_loss:.4f}, acc={tr_acc:.4f} | "
          f"Val loss={va_loss:.4f}, acc={va_acc:.4f}")

    if va_acc > best_val_acc:
        best_val_acc = va_acc
        best_state = {k: v.cpu() for k, v in model.state_dict().items()}

if best_state is not None:
    model.load_state_dict(best_state)
    model = model.to(device)

print("Best val acc:", best_val_acc)

```

```

Epoch 1/10 | Train loss=0.8239, acc=0.7554 | Val loss=0.3957, acc=0.8817
Epoch 2/10 | Train loss=0.5305, acc=0.8326 | Val loss=0.3148, acc=0.9033
Epoch 3/10 | Train loss=0.4863, acc=0.8439 | Val loss=0.3107, acc=0.9028
Epoch 4/10 | Train loss=0.4674, acc=0.8491 | Val loss=0.2777, acc=0.9136
Epoch 5/10 | Train loss=0.4524, acc=0.8536 | Val loss=0.2988, acc=0.9048
Epoch 6/10 | Train loss=0.4460, acc=0.8553 | Val loss=0.2949, acc=0.9063
Epoch 7/10 | Train loss=0.4397, acc=0.8573 | Val loss=0.2955, acc=0.9097
Epoch 8/10 | Train loss=0.4334, acc=0.8598 | Val loss=0.2877, acc=0.9112

```

```

Epoch 9/10 | Train loss=0.4326, acc=0.8582 | Val loss=0.2906, acc=0.9024
Epoch 10/10 | Train loss=0.4291, acc=0.8612 | Val loss=0.3078, acc=0.9048
Best val acc: 0.9136408243375859

```

Stage 2: Unfreeze layer 4 and fc

```
[13]: for name, p in model.named_parameters():
    p.requires_grad = name.startswith("layer4") or name.startswith("fc")

optimizer = optim.AdamW([
    {"params": model.layer4.parameters(), "lr": 1e-4},
    {"params": model.fc.parameters(), "lr": 5e-4},
], weight_decay=1e-4)

scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=5)

num_epochs_ft = 5
best_val_acc_ft = 0.0
best_state_ft = None

for epoch in range(num_epochs_ft):
    tr_loss, tr_acc = train_one_epoch(model, train_loader, optimizer, criterion, device)
    va_loss, va_acc = eval_one_epoch(model, val_loader, criterion, device)

    print(f"[FT] Epoch {epoch+1}/{num_epochs_ft} | "
          f"Train loss={tr_loss:.4f}, acc={tr_acc:.4f} | "
          f"Val loss={va_loss:.4f}, acc={va_acc:.4f}", flush=True)

    scheduler.step()

    if va_acc > best_val_acc_ft:
        best_val_acc_ft = va_acc
        best_state_ft = {k: v.cpu() for k, v in model.state_dict().items()}

if best_state_ft is not None:
    model.load_state_dict(best_state_ft)
    model = model.to(device)

print("Best FT val acc:", best_val_acc_ft)
```

```

[FT] Epoch 1/5 | Train loss=0.4423, acc=0.8595 | Val loss=0.2472, acc=0.9249
[FT] Epoch 2/5 | Train loss=0.3057, acc=0.9002 | Val loss=0.2236, acc=0.9338
[FT] Epoch 3/5 | Train loss=0.2463, acc=0.9201 | Val loss=0.1948, acc=0.9460
[FT] Epoch 4/5 | Train loss=0.1929, acc=0.9364 | Val loss=0.1782, acc=0.9460
[FT] Epoch 5/5 | Train loss=0.1688, acc=0.9443 | Val loss=0.1680, acc=0.9509
Best FT val acc: 0.9509322865554465

```

5 4. Fine tuned Mobile Net V3

```
[14]: mb_weights = models.MobileNet_V3_Large_Weights.IMGNET1K_V1
mobile_model = models.mobilenet_v3_large(weights=mb_weights)
mobile_model.classifier[-1] = nn.Linear(mobile_model.classifier[-1].
    ↪in_features, num_fine_classes)

for p in mobile_model.features.parameters():
    p.requires_grad = False

mobile_model = mobile_model.to(device)
mobile_criterion = nn.CrossEntropyLoss()
mobile_optimizer = optim.Adam(mobile_model.classifier.parameters(), lr=1e-3)
mobile_epochs = 10
mobile_best_acc = 0.0
mobile_best_state = None

for epoch in range(mobile_epochs):
    tr_loss, tr_acc = train_one_epoch(mobile_model, train_loader, ↪
        ↪mobile_optimizer, mobile_criterion, device)
    va_loss, va_acc = eval_one_epoch(mobile_model, val_loader, ↪
        ↪mobile_criterion, device)
    print(f"[MobileNet] Epoch {epoch+1}/{mobile_epochs} | Train loss={tr_loss:.4f}, acc={tr_acc:.4f} | Val loss={va_loss:.4f}, acc={va_acc:.4f}")
    if va_acc > mobile_best_acc:
        mobile_best_acc = va_acc
        mobile_best_state = {k: v.cpu() for k, v in mobile_model.state_dict().items()}

if mobile_best_state is not None:
    mobile_model.load_state_dict(mobile_best_state)
    mobile_model = mobile_model.to(device)
print("[MobileNet] Best val acc:", mobile_best_acc)

# fine-tune: unfreeze backbone with smaller LR for a few epochs
for p in mobile_model.features.parameters():
    p.requires_grad = True

mobile_ft_optimizer = optim.AdamW([
    {"params": mobile_model.features.parameters(), "lr": 3e-5, "weight_decay": ↪
        ↪1e-4},
    {"params": mobile_model.classifier.parameters(), "lr": 3e-4, "weight_decay": ↪
        ↪ 1e-4},
])
mobile_ft_epochs = 5
mobile_ft_best_acc = 0.0
mobile_ft_best_state = None
```

```

for epoch in range(mobile_ft_epochs):
    tr_loss, tr_acc = train_one_epoch(mobile_model, train_loader,
                                      mobile_ft_optimizer, mobile_criterion, device)
    va_loss, va_acc = eval_one_epoch(mobile_model, val_loader,
                                      mobile_criterion, device)
    print(f"[MobileNet-FT] Epoch {epoch+1}/{mobile_ft_epochs} | Train loss={tr_loss:.4f}, acc={tr_acc:.4f} | Val loss={va_loss:.4f}, acc={va_acc:.4f}")
    if va_acc > mobile_ft_best_acc:
        mobile_ft_best_acc = va_acc
        mobile_ft_best_state = {k: v.cpu() for k, v in mobile_model.state_dict().items()}

if mobile_ft_best_state is not None:
    mobile_model.load_state_dict(mobile_ft_best_state)
    mobile_model = mobile_model.to(device)
print("[MobileNet-FT] Best val acc:", mobile_ft_best_acc)

```

Downloading:

```

"https://download.pytorch.org/models/mobilenet_v3_large-8738ca79.pth" to
/root/.cache/torch/hub/checkpoints/mobilenet_v3_large-8738ca79.pth

100% | 21.1M/21.1M [00:00<00:00, 70.7MB/s]

[MobileNet] Epoch 1/10 | Train loss=0.6210, acc=0.8090 | Val loss=0.3032,
acc=0.9053
[MobileNet] Epoch 2/10 | Train loss=0.4433, acc=0.8573 | Val loss=0.2660,
acc=0.9171
[MobileNet] Epoch 3/10 | Train loss=0.4037, acc=0.8709 | Val loss=0.2370,
acc=0.9303
[MobileNet] Epoch 4/10 | Train loss=0.3640, acc=0.8824 | Val loss=0.2605,
acc=0.9190
[MobileNet] Epoch 5/10 | Train loss=0.3515, acc=0.8888 | Val loss=0.2331,
acc=0.9323
[MobileNet] Epoch 6/10 | Train loss=0.3385, acc=0.8909 | Val loss=0.2535,
acc=0.9195
[MobileNet] Epoch 7/10 | Train loss=0.3201, acc=0.8978 | Val loss=0.2629,
acc=0.9195
[MobileNet] Epoch 8/10 | Train loss=0.3129, acc=0.8996 | Val loss=0.2550,
acc=0.9220
[MobileNet] Epoch 9/10 | Train loss=0.3127, acc=0.9023 | Val loss=0.2167,
acc=0.9342
[MobileNet] Epoch 10/10 | Train loss=0.2988, acc=0.9058 | Val loss=0.2332,
acc=0.9347
[MobileNet] Best val acc: 0.9347399411187438
[MobileNet-FT] Epoch 1/5 | Train loss=0.2264, acc=0.9279 | Val loss=0.1811,
acc=0.9534

```

```
[MobileNet-FT] Epoch 2/5 | Train loss=0.1710, acc=0.9435 | Val loss=0.1731,
acc=0.9509
[MobileNet-FT] Epoch 3/5 | Train loss=0.1692, acc=0.9464 | Val loss=0.1521,
acc=0.9588
[MobileNet-FT] Epoch 4/5 | Train loss=0.1400, acc=0.9555 | Val loss=0.1504,
acc=0.9598
[MobileNet-FT] Epoch 5/5 | Train loss=0.1349, acc=0.9561 | Val loss=0.1582,
acc=0.9578
[MobileNet-FT] Best val acc: 0.9597644749754661
```

6 5. Strategies evaluation and comparison

```
[15]: import numpy as np
import torch
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
from matplotlib.colors import PowerNorm, LogNorm
import matplotlib.pyplot as plt
import numpy as np

@torch.no_grad()
def collect_fine_preds(model, loader, device):
    model.eval()
    y_true, y_pred = [], []

    for images, fine_labels, _ in loader:
        images = images.to(device)
        outputs = model(images)
        pred = outputs.argmax(dim=1).cpu().numpy()

        y_true.extend(fine_labels.numpy())
        y_pred.extend(pred)

    return np.array(y_true), np.array(y_pred)

def _plot_confusion_matrix(cm, class_names, title, normalize=False,
                           cmap="Blues", use_log_for_counts=True, ↴
                           vmax_percentile=99):
    cm_plot = cm.astype(np.float64)

    if normalize:
        row_sums = cm_plot.sum(axis=1, keepdims=True)
        cm_plot = np.divide(cm_plot, row_sums, out=np.zeros_like(cm_plot), ↴
                           where=row_sums != 0) * 100.0
```

```

n = len(class_names)

fig_w = min(18, max(10, 0.65 * n))
fig_h = min(18, max(7, 0.60 * n))
plt.figure(figsize=(fig_w, fig_h))

if normalize:

    vmax = 100.0
    norm = PowerNorm(gamma=0.5, vmin=0.0, vmax=vmax)
    im = plt.imshow(cm_plot, interpolation="nearest", cmap=cmap, norm=norm)
else:

    nonzero = cm_plot[cm_plot > 0]
    vmax = np.percentile(nonzero, vmax_percentile) if nonzero.size else 1.0

    if use_log_for_counts:

        masked = np.ma.masked_where(cm_plot == 0, cm_plot)
        norm = LogNorm(vmin=1, vmax=max(1, vmax))
        im = plt.imshow(masked, interpolation="nearest", cmap=cmap, norm=norm)
    else:

        norm = PowerNorm(gamma=0.5, vmin=0.0, vmax=vmax)
        im = plt.imshow(cm_plot, interpolation="nearest", cmap=cmap, norm=norm)

plt.title(title)
plt.colorbar(im, fraction=0.046, pad=0.04)

tick_marks = np.arange(n)
plt.xticks(tick_marks, class_names, rotation=45, ha="right", fontsize=9)
plt.yticks(tick_marks, class_names, fontsize=9)

display_max = np.nanmax(cm_plot) if normalize else (np.
    nanmax(cm_plot[cm_plot > 0]) if np.any(cm_plot > 0) else 1.0)
thresh = display_max * 0.5

for i in range(n):
    for j in range(n):
        val = cm_plot[i, j]
        if normalize:
            text = f"{val:.1f}"
            show = (val > 0)

```

```

    else:
        text = str(int(cm[i, j]))
        show = (cm[i, j] > 0)

    if show:
        plt.text(j, i, text,
                  ha="center", va="center",
                  fontsize=8,
                  color="white" if val > thresh else "black")

plt.ylabel("True label")
plt.xlabel("Predicted label")
plt.tight_layout()
plt.show()

def evaluate_report(y_fine_true, y_fine_pred, fine_idx_to_name, fine_to_big, ↴
                     big_idx_to_name, set_name=""):

    fine_names = [fine_idx_to_name[i] for i in range(len(fine_idx_to_name))]

    print(f"\n==== {set_name} Fine-class Report ===")
    print(classification_report(y_fine_true, y_fine_pred, ↴
                                target_names=fine_names, zero_division=0))

    cm_fine = confusion_matrix(y_fine_true, y_fine_pred, labels=np.
                               arange(len(fine_names)))
    _plot_confusion_matrix(cm_fine, fine_names, title=f"{set_name} Fine ↴
                           Confusion Matrix (Counts)", normalize=False)
    _plot_confusion_matrix(cm_fine, fine_names, title=f"{set_name} Fine ↴
                           Confusion Matrix (Row %)", normalize=True)

    y_big_true = np.array([fine_to_big[int(i)] for i in y_fine_true])
    y_big_pred = np.array([fine_to_big[int(i)] for i in y_fine_pred])

    big_names = [big_idx_to_name[i] for i in range(len(big_idx_to_name))]

    print(f"\n==== {set_name} Big-class Report ===")
    print(classification_report(y_big_true, y_big_pred, target_names=big_names, ↴
                                zero_division=0))

    cm_big = confusion_matrix(y_big_true, y_big_pred, labels=np.
                               arange(len(big_names)))
    _plot_confusion_matrix(cm_big, big_names, title=f"{set_name} Big Confusion ↴
                           Matrix (Counts)", normalize=False)

```

```

    _plot_confusion_matrix(cm_big, big_names, title=f"{set_name} Big Confusion_
    ↵Matrix (Row %)", normalize=True)

print("\n##### SVM VAL SET EVAL #####")
evaluate_report(svm_val_labels, svm_val_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="SVM VAL")

print("\n##### SVM TEST SET EVAL #####")
evaluate_report(svm_test_labels, svm_test_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="SVM TEST")

print("\n##### CNN VAL SET EVAL #####")
y_cnn_val_true, y_cnn_val_pred = collect_fine_preds(cnn_model, val_loader, ↵
    ↵device)
evaluate_report(y_cnn_val_true, y_cnn_val_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="CNN VAL")

print("\n##### CNN TEST SET EVAL #####")
y_cnn_test_true, y_cnn_test_pred = collect_fine_preds(cnn_model, test_loader, ↵
    ↵device)
evaluate_report(y_cnn_test_true, y_cnn_test_pred, fine_idx_to_name, fine_to_big, ↵
    ↵fine_to_big, big_idx_to_name, set_name="CNN TEST")

y_val_true, y_val_pred = collect_fine_preds(model, val_loader, device)
print("\n##### ResNet VAL SET EVAL #####")
evaluate_report(y_val_true, y_val_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="VAL")

y_test_true, y_test_pred = collect_fine_preds(model, test_loader, device)
print("\n##### ResNet TEST SET EVAL #####")
evaluate_report(y_test_true, y_test_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="TEST")

print("\n##### MobileNet VAL SET EVAL #####")
y_mb_val_true, y_mb_val_pred = collect_fine_preds(mobile_model, val_loader, ↵
    ↵device)
evaluate_report(y_mb_val_true, y_mb_val_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="MobileNet VAL")

print("\n##### MobileNet TEST SET EVAL #####")
y_mb_test_true, y_mb_test_pred = collect_fine_preds(mobile_model, test_loader, ↵
    ↵device)
evaluate_report(y_mb_test_true, y_mb_test_pred, fine_idx_to_name, fine_to_big, ↵
    ↵big_idx_to_name, set_name="MobileNet TEST")

```

Output hidden; open in <https://colab.research.google.com> to view.

```
[ ]: # from https://gist.github.com/jonathanagustin/b67b97ef12c53a8dec27b343dca4abba
# install can take a minute

import os
# @title Convert Notebook to PDF. Save Notebook to given directory
NOTEBOOKS_DIR = "/content/drive/My Drive/CS441/Final" # @param {type:"string"}
NOTEBOOK_NAME = "Imagenet.ipynb" # @param {type:"string"}
#-----#
from google.colab import drive
drive.mount("/content/drive/", force_remount=True)
NOTEBOOK_PATH = f"{NOTEBOOKS_DIR}/{NOTEBOOK_NAME}"
assert os.path.exists(NOTEBOOK_PATH), f"NOTEBOOK NOT FOUND: {NOTEBOOK_PATH}"
!apt install -y texlive-xetex texlive-fonts-recommended texlive-plain-generic > /dev/null 2>&1
!jupyter nbconvert "$NOTEBOOK_PATH" --to pdf > /dev/null 2>&1
NOTEBOOK_PDF = NOTEBOOK_PATH.rsplit('.', 1)[0] + '.pdf'
assert os.path.exists(NOTEBOOK_PDF), f"ERROR MAKING PDF: {NOTEBOOK_PDF}"
print(f"PDF CREATED: {NOTEBOOK_PDF}")
```

Mounted at /content/drive/